

Bank Marketing Campaign Data

Introduction:

The prediction model uses direct marketing campaign data from a banking institution. Phone calls were used in the marketing campaigns to contact customers in order to determine whether the customer would subscribe to a term deposit.

The objective of this project is to create a model that will predict whether or not a customer will sign up for the product (bank term deposit) based on the data at hand.

Data Specifications:

1. Age - Age of the customer.
2. Job - Profession of the customer.
3. Marital Status - Relationship status of the customer.
4. Education - Education level of the customer.
5. Default - / Having a previously broken credit.
6. Housing - Whether customer has a home loan.
7. Loan - Whether the customer has a personal loan.
8. Contact - Whether the customer was contacted on his home or mobile phone.
9. Month - The last month when the customer was contacted.
10. Day - The day when the customer was contacted.
11. Duration - The total time spent on talking to the customer.
12. Campaign - The number of contacts reaching the customer during the current campaign (including the last contact).
13. Pdays: The number of days since the previous campaign, if reached (-1 if it was never reached before).
14. Previous: The number of contacts that reached the customer before this campaign.
15. Poutcome: Previous campaign result, whether successful or not.
16. Deposit: Whether customer has a term deposit or not.

Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Loading data

```
In [2]: data = pd.read_csv('C:/Data Science Projects/Product Uptake Prediction/bank.csv')
data.head()
```

```
Out[2]:
```

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campa
0	59	admin.	married	secondary	no	2343	yes	no	unknown	5	may	1042	

1	56	admin.	married	secondary	no	45	no	no	unknown	5	may	1467
2	41	technician	married	secondary	no	1270	yes	no	unknown	5	may	1389
3	55	services	married	secondary	no	2476	yes	no	unknown	5	may	579
4	54	admin.	married	tertiary	no	184	no	no	unknown	5	may	673

Checking dataframe info

number of columns, column labels, column data types, memory usage, range index, and the number of cells in each column (non-null values)

In [3]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11162 entries, 0 to 11161
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         11162 non-null  int64
 1   job         11162 non-null  object
 2   marital     11162 non-null  object
 3   education   11162 non-null  object
 4   default     11162 non-null  object
 5   balance     11162 non-null  int64
 6   housing     11162 non-null  object
 7   loan        11162 non-null  object
 8   contact     11162 non-null  object
 9   day         11162 non-null  int64
10  month       11162 non-null  object
11  duration    11162 non-null  int64
12  campaign    11162 non-null  int64
13  pdays      11162 non-null  int64
14  previous    11162 non-null  int64
15  poutcome    11162 non-null  object
16  deposit     11162 non-null  object
dtypes: int64(7), object(10)
memory usage: 1.4+ MB
```

Checking missing values

(In this case 'NA's as there are no null values as observed above).

In [4]: `data.isna().sum()`

```
Out[4]: age         0
job         0
marital     0
education   0
default     0
balance     0
housing     0
loan        0
contact     0
day         0
month       0
duration    0
campaign    0
pdays      0
previous    0
poutcome    0
```

```
deposit      0  
dtype: int64
```

Outlier detection

```
In [5]: from collections import Counter  
  
def detect_outliers(data, features):  
  
    outlier_indices = []  
  
    for f in features:  
  
        # 1st quartile  
        Q1 = np.percentile(data[f], 25)  
  
        # 3rd quartile  
        Q3 = np.percentile(data[f], 75)  
  
        # IQR  
        IQR = Q3 - Q1  
  
        # Outlier size  
        outlier_size = IQR * 1.5  
  
        # detect outliers indices  
        outlier_list = data[(data[f] < Q1 - outlier_size) | (data[f] > Q3 + outlier_size)]  
  
        # store outlier indices  
        outlier_indices.extend(outlier_list)  
  
    outlier_indices = Counter(outlier_indices)  
  
    multiple_outliers = list(i for i, v in outlier_indices.items() if v > 2)  
  
    return multiple_outliers
```

Identify outliers based on the features (duration, campaign, and previous) directly linked to the marketing campaign.

```
In [6]: data.loc[detect_outliers(data, ['duration', 'campaign', 'previous'])]
```

```
Out[6]:   age  job  marital  education  default  balance  housing  loan  contact  day  month  duration  campaign  pda
```

No outlier detected.

NB\: Unless outliers are caused by data entry errors, usually outliers should not be removed as they represent the natural occurrence of the data.

Descriptive statistics

Numerical variables

```
In [7]: data.describe()
```

```
Out[7]:
```

	age	balance	day	duration	campaign	pdays	previous
count	11162.000000	11162.000000	11162.000000	11162.000000	11162.000000	11162.000000	11162.000000
mean	41.231948	1528.538524	15.658036	371.993818	2.508421	51.330407	0.832557

std	11.913369	3225.413326	8.420740	347.128386	2.722077	108.758282	2.292007
min	18.000000	-6847.000000	1.000000	2.000000	1.000000	-1.000000	0.000000
25%	32.000000	122.000000	8.000000	138.000000	1.000000	-1.000000	0.000000
50%	39.000000	550.000000	15.000000	255.000000	2.000000	-1.000000	0.000000
75%	49.000000	1708.000000	22.000000	496.000000	3.000000	20.750000	1.000000
max	95.000000	81204.000000	31.000000	3881.000000	63.000000	854.000000	58.000000

```
In [8]: num_features = ['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']
```

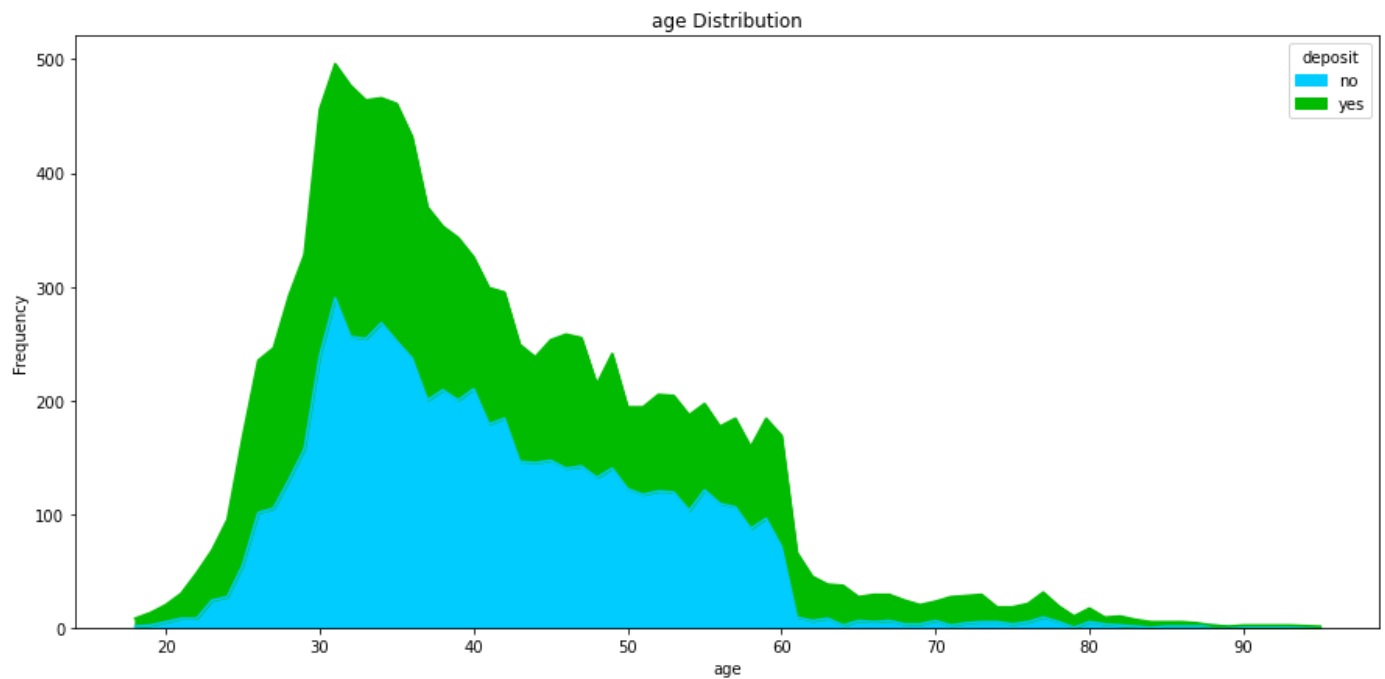
```
In [9]: def area_plot(feature):
    var = data[feature]


    pd.crosstab(var, data.deposit).plot(kind="area", figsize=(15, 7), color=['#00ccff', '#00b
    plt.title("{} Distribution".format(feature))
    plt.xlabel("{}".format(feature))
    plt.ylabel('Frequency')
    plt.show()
```

```
In [10]: # Considering only age and balance
```

```
num_features1 = ['age', 'balance']
```

```
for f in num_features1:
    area_plot(f)
```





Age:

From the age distribution graph above, we can note that the highest number of customers with term deposits range between the ages 30 - 35 years. Also, within the same age range, there are many customers who have not subscribed to a term deposit, hence there is still a huge opportunity to sell this product others customers within this age range.

Balance

From the balance distribution graph above, we note that majority of the customers who have subscribed to term deposits have a nearly zero balance. we may observe this more clearly by grouping the customers based on their balances.

```
In [11]: # Categorizing the customers by balance

data1 = data.copy()

data1.loc[data1['balance'] >= 10000, 'bal_cat'] = "High_bal"
data1.loc[(data1['balance'] > 2000) & (data1['balance'] < 10000), 'bal_cat'] = "Average_b
data1.loc[(data1['balance'] >= 0) & (data1['balance'] < 2000), 'bal_cat'] = "Low_bal"
data1.loc[data1['balance'] < 0, 'bal_cat'] = "Negative_bal"

table_bal_cat = pd.crosstab(data1['deposit'], data1['bal_cat']).apply(lambda x: x/x.sum(
table_bal_cat
```

```
Out[11]: bal_cat  Average_bal  High_bal  Low_bal  Negative_bal
```

deposit


no	42.0	41.0	54.0	69.0
yes	58.0	59.0	46.0	31.0

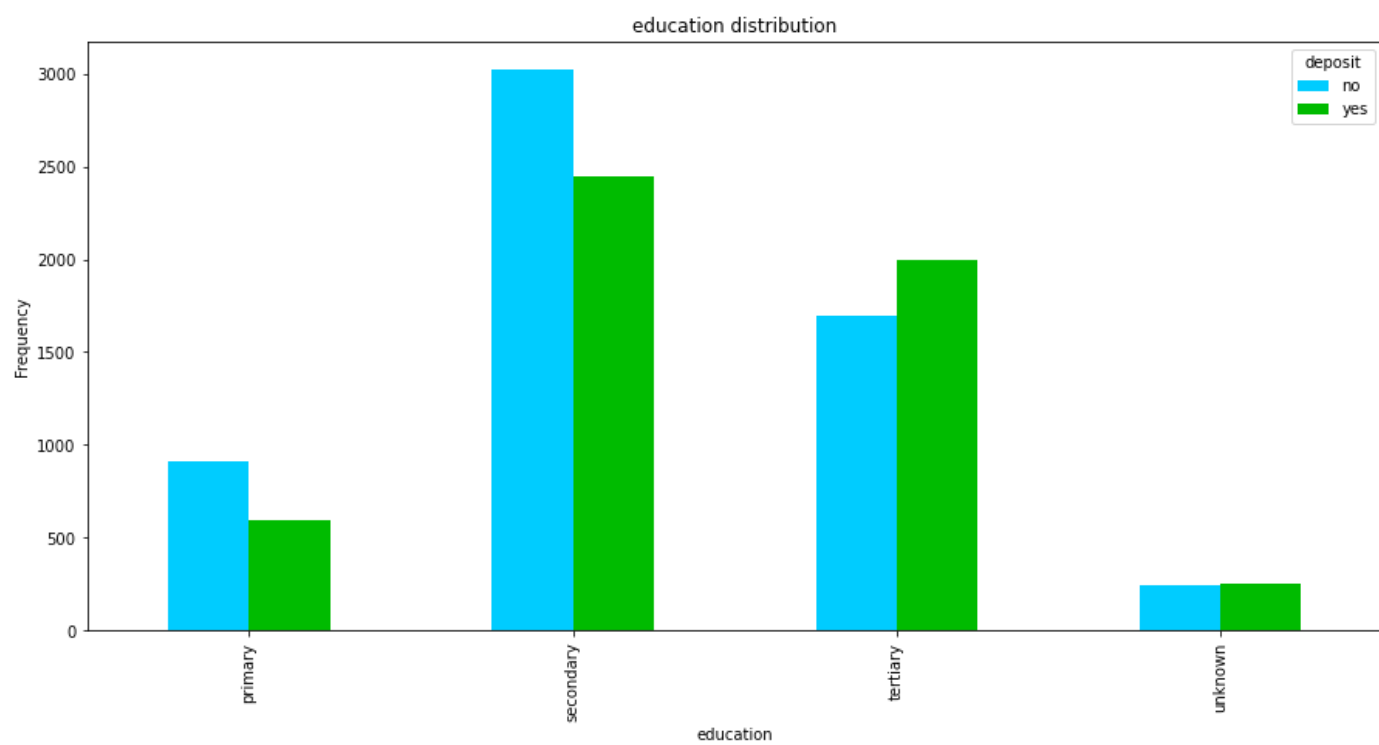
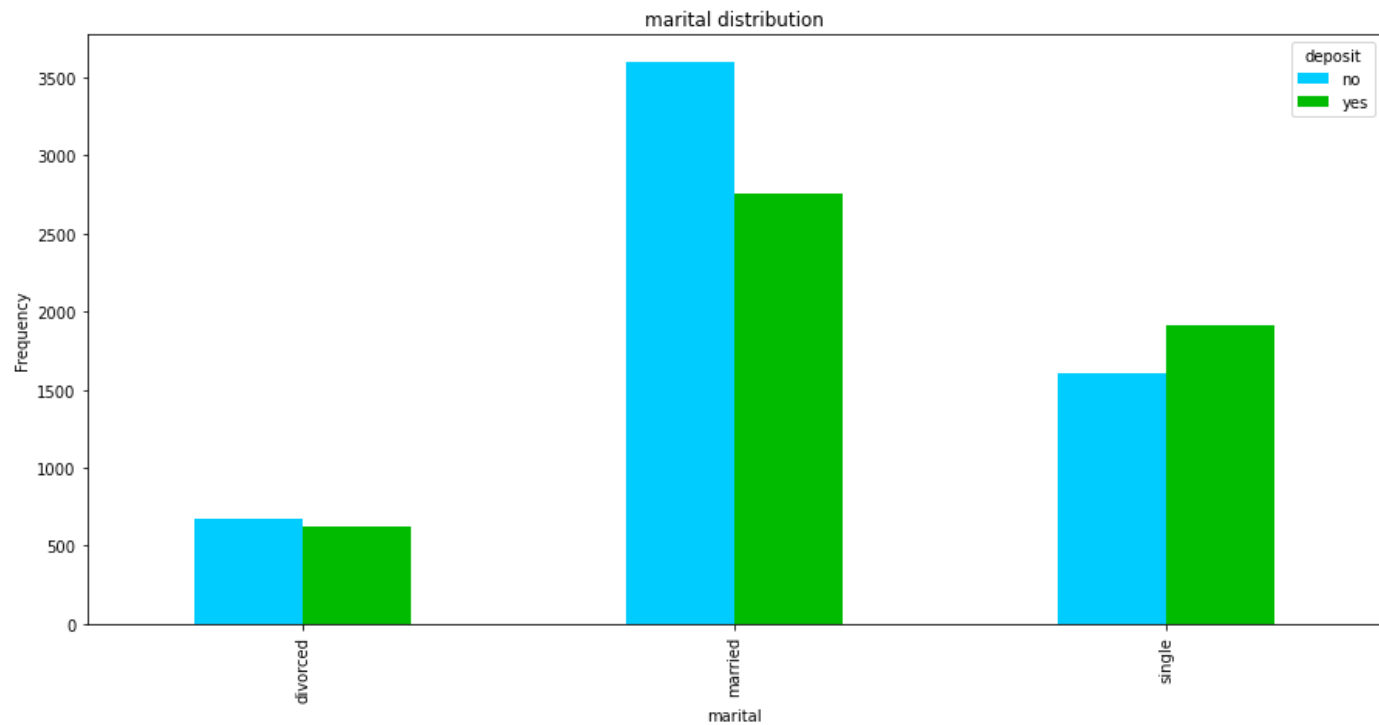
From table above we can note that nearly 60% of those customers having average and high balances (over 2,000) have a term deposit. Important to note is that quite a significant proportion of customers with low and negative balances also have a term deposit. This indicates that the marketing team can target customers with any balance amount.

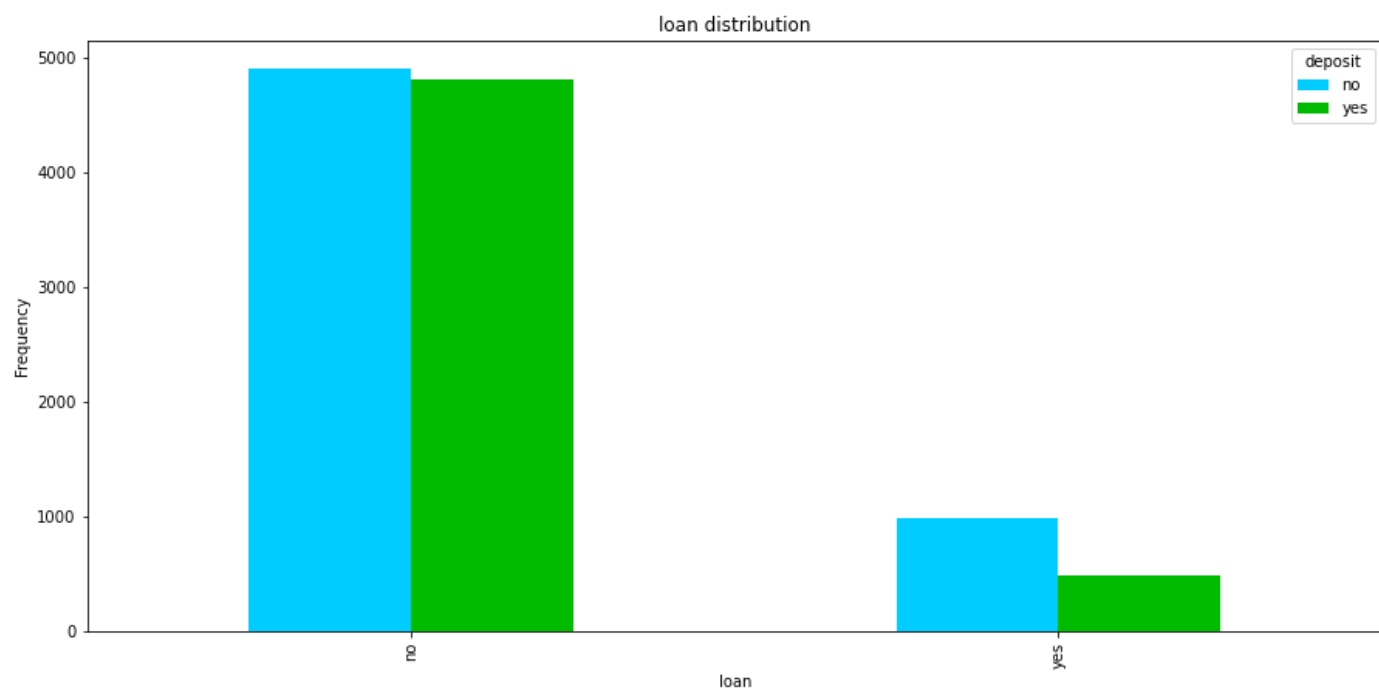
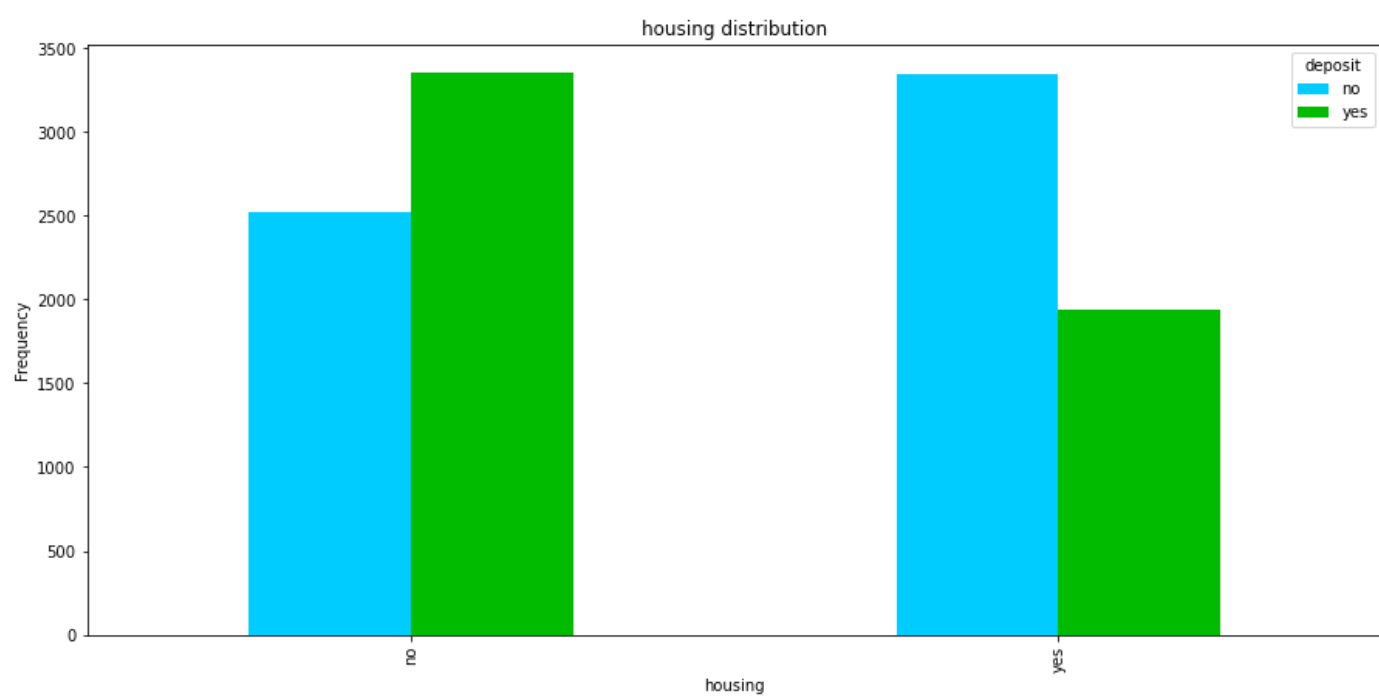
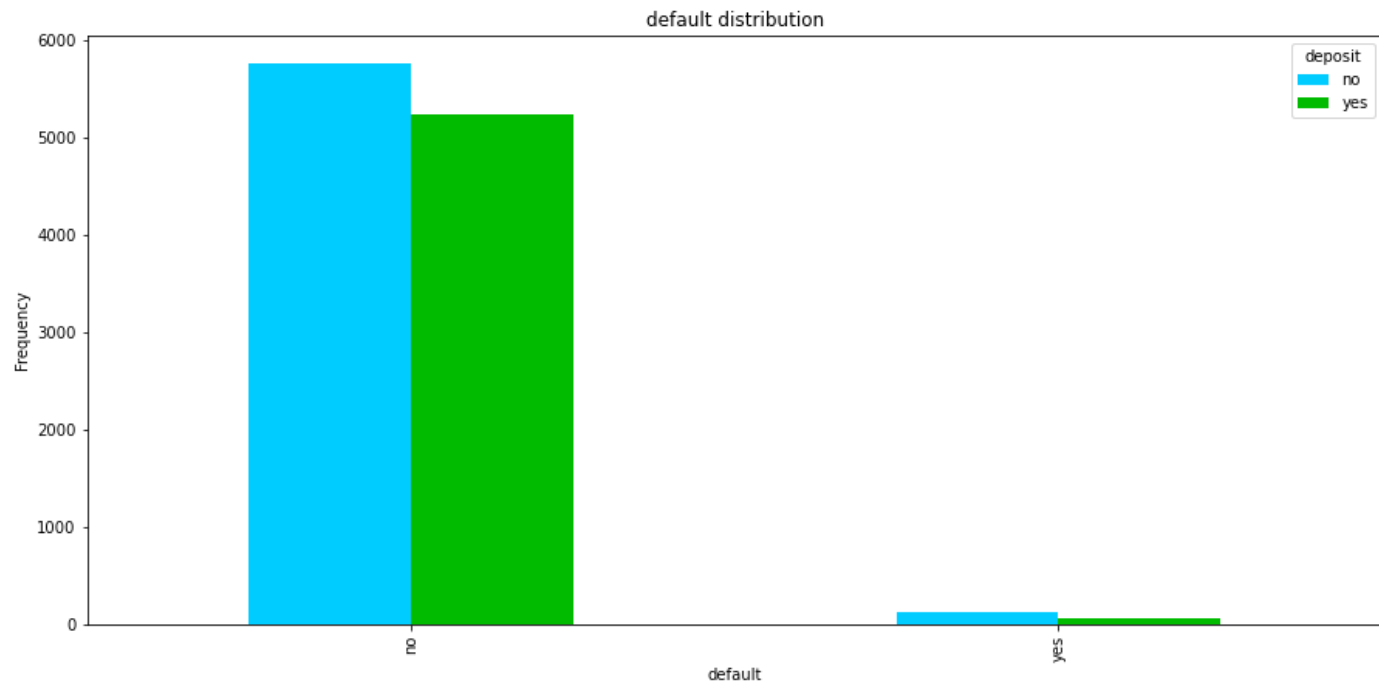
Categorical variables



```
In [12]: def bar_plot(feature):  
    var = data[feature]  
  
    pd.crosstab(var,data.deposit).plot(kind="bar",figsize=(15,7),color=['#00ccff','#00bb  
    plt.title("{} distribution".format(feature))  
    plt.xlabel("{}".format(feature))  
    plt.ylabel('Frequency')  
    plt.show()
```

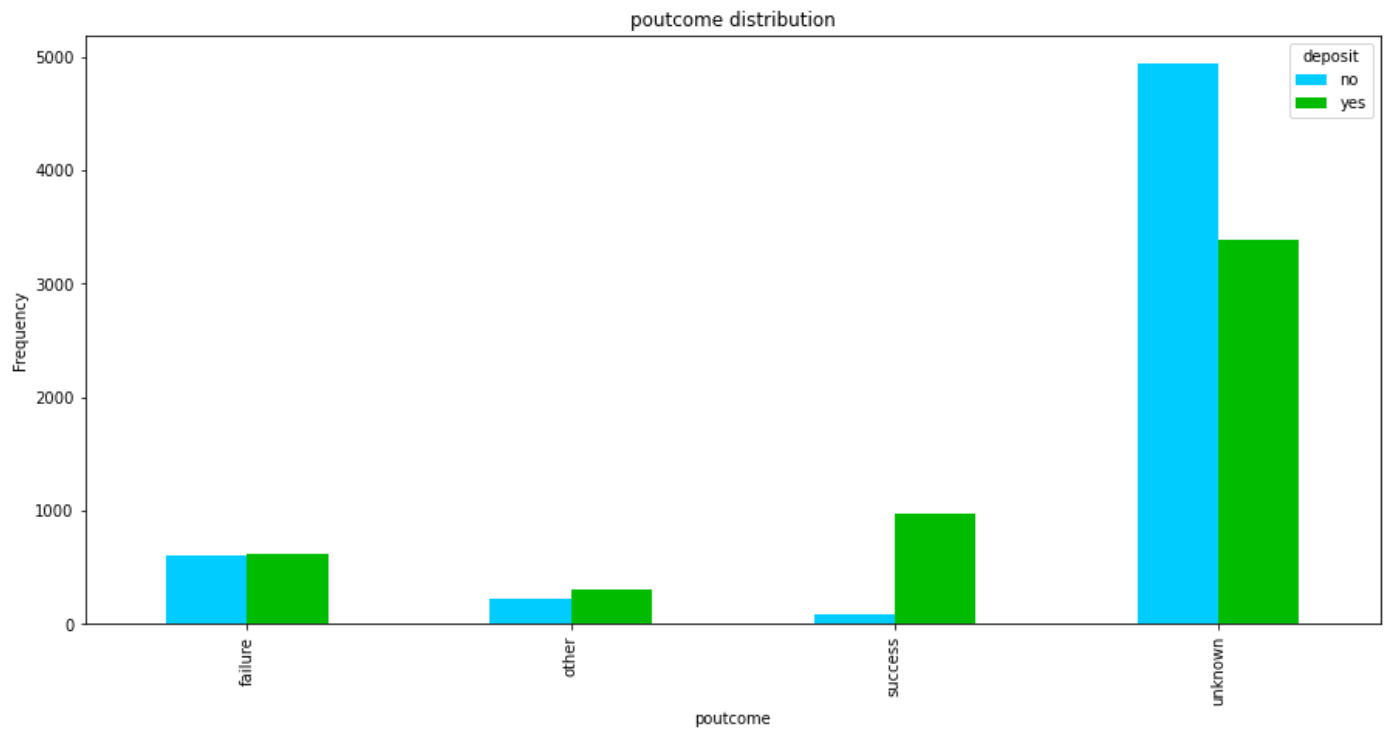
```
In [13]: cat_features = ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', '  
  
for f in cat_features:  
    bar_plot(f)
```











The graphs above can provide an overview of the characteristics of customers with high number of term deposits. For example, from the first graph, we can note that customers with a management job have the highest number of term deposits and so on. For a detailed summary of the exact numbers and subscription proportions of each variable see below.

```
In [14]: for f in cat_features:
          print(pd.crosstab(data['deposit'], data[f]).apply(lambda x: x/x.sum() * 100).round())
          print(data.groupby(f).size().sort_values(ascending=False), '\n')
```

job	admin.	blue-collar	entrepreneur	housemaid	management	retired	\
deposit							
no	53.0	64.0	62.0	60.0	49.0	34.0	
yes	47.0	36.0	38.0	40.0	51.0	66.0	

job	self-employed	services	student	technician	unemployed	unknown
deposit						
no	54.0	60.0	25.0	54.0	43.0	51.0
yes	46.0	40.0	75.0	46.0	57.0	49.0

```
job
management      2566
blue-collar      1944
technician       1823
admin.           1334
services         923
retired          778
self-employed    405
student          360
unemployed       357
entrepreneur     328
housemaid        274
unknown          70
dtype: int64
```

marital	divorced	married	single
deposit			
no	52.0	57.0	46.0
yes	48.0	43.0	54.0

marital
married 6351
single 3518
divorced 1293
dtype: int64

education	primary	secondary	tertiary	unknown
deposit				
no	61.0	55.0	46.0	49.0
yes	39.0	45.0	54.0	51.0

education
secondary 5476
tertiary 3689
primary 1500
unknown 497
dtype: int64

default	no	yes
deposit		
no	52.0	69.0
yes	48.0	31.0

default
no 10994
yes 168
dtype: int64

housing	no	yes
deposit		
no	43.0	63.0
yes	57.0	37.0

housing
no 5881
yes 5281
dtype: int64

loan	no	yes
deposit		
no	50.0	67.0
yes	50.0	33.0

loan
no 9702
yes 1460
dtype: int64

contact	cellular	telephone	unknown
deposit			
no	46.0	50.0	77.0
yes	54.0	50.0	23.0

contact
cellular 8042
unknown 2346
telephone 774
dtype: int64

month	apr	aug	dec	feb	jan	jul	jun	mar	may	nov	oct	\
deposit												
no	37.0	55.0	9.0	43.0	59.0	59.0	55.0	10.0	67.0	57.0	18.0	
yes	63.0	45.0	91.0	57.0	41.0	41.0	45.0	90.0	33.0	43.0	82.0	

month sep

```

deposit
no      16.0
yes     84.0

month
may      2824
aug      1519
jul      1514
jun      1222
nov       943
apr       923
feb       776
oct       392
jan       344
sep       319
mar       276
dec       110
dtype: int64

poutcome  failure  other  success  unknown
deposit
no          50.0   43.0     9.0    59.0
yes          50.0   57.0    91.0    41.0

poutcome
unknown     8326
failure     1228
success     1071
other        537
dtype: int64

```

Correlations

```

In [15]: fig, ax = plt.subplots(figsize=(15,7)) # Sample figsize in inches
sns.heatmap(data.corr(), annot=True, linewidths=.5, ax=ax)

```

```

Out[15]: <AxesSubplot:>

```



Majority of the variables have weak positive or negative correlations except the relationship between the

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11162 entries, 0 to 11161
Data columns (total 52 columns):
#   Column                                Non-Null Count  Dtype

```

0	age	11162	non-null	int64
1	balance	11162	non-null	int64
2	day	11162	non-null	int64
3	duration	11162	non-null	int64
4	campaign	11162	non-null	int64
5	pdays	11162	non-null	int64
6	previous	11162	non-null	int64
7	deposit	11162	non-null	int64
8	job_admin.	11162	non-null	uint8
9	job_blue-collar	11162	non-null	uint8
10	job_entrepreneur	11162	non-null	uint8
11	job_housemaid	11162	non-null	uint8
12	job_management	11162	non-null	uint8
13	job_retired	11162	non-null	uint8
14	job_self-employed	11162	non-null	uint8
15	job_services	11162	non-null	uint8
16	job_student	11162	non-null	uint8
17	job_technician	11162	non-null	uint8
18	job_unemployed	11162	non-null	uint8
19	job_unknown	11162	non-null	uint8
20	marital_divorced	11162	non-null	uint8
21	marital_married	11162	non-null	uint8
22	marital_single	11162	non-null	uint8
23	education_primary	11162	non-null	uint8
24	education_secondary	11162	non-null	uint8
25	education_tertiary	11162	non-null	uint8
26	education_unknown	11162	non-null	uint8
27	default_no	11162	non-null	uint8
28	default_yes	11162	non-null	uint8
29	housing_no	11162	non-null	uint8
30	housing_yes	11162	non-null	uint8
31	loan_no	11162	non-null	uint8
32	loan_yes	11162	non-null	uint8
33	contact_cellular	11162	non-null	uint8
34	contact_telephone	11162	non-null	uint8
35	contact_unknown	11162	non-null	uint8
36	month_apr	11162	non-null	uint8
37	month_aug	11162	non-null	uint8
38	month_dec	11162	non-null	uint8
39	month_feb	11162	non-null	uint8
40	month_jan	11162	non-null	uint8
41	month_jul	11162	non-null	uint8
42	month_jun	11162	non-null	uint8
43	month_mar	11162	non-null	uint8
44	month_may	11162	non-null	uint8
45	month_nov	11162	non-null	uint8
46	month_oct	11162	non-null	uint8
47	month_sep	11162	non-null	uint8
48	poutcome_failure	11162	non-null	uint8
49	poutcome_other	11162	non-null	uint8
50	poutcome_success	11162	non-null	uint8
51	poutcome_unknown	11162	non-null	uint8

dtypes: int64(8), uint8(44)
memory usage: 1.1 MB

Splitting the dataset

```
In [19]: # Splitting dataset into independent (X) and dependent (y) variables

X=data_main.drop(['deposit'], axis=1)

y=data_main['deposit']
```

Test Train Split

```
In [20]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

Feature scaling

```
In [21]: from sklearn.preprocessing import StandardScaler

sc = StandardScaler()

X_train = sc.fit_transform(X_train)

X_test = sc.transform(X_test)
```

Modeling

```
In [40]: from sklearn.model_selection import train_test_split
from sklearn.metrics import cohen_kappa_score
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import f1_score
```

Logistic Regression

```
In [41]: from sklearn.linear_model import LogisticRegression

lr_model = LogisticRegression()

lr_model.fit(X_train,y_train)

y_pred = lr_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels = ['0', '1'])


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("Logistic Regression accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.81	0.83	0.82	1166
1	0.81	0.78	0.80	1067
accuracy			0.81	2233
macro avg	0.81	0.81	0.81	2233
weighted avg	0.81	0.81	0.81	2233

Logistic Regression accuracy: 80.83

F1-Score: 79.56

Cohen Kappa score: 61.53

Naive Bayes

```
In [27]: from sklearn.naive_bayes import GaussianNB

nb_model = GaussianNB()

nb_model.fit(X_train,y_train)

y_pred = nb_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("Naive Bayes accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```

	precision	recall	f1-score	support
0	0.69	0.87	0.77	1166
1	0.80	0.57	0.67	1067
accuracy			0.73	2233
macro avg	0.75	0.72	0.72	2233
weighted avg	0.74	0.73	0.72	2233

Naive Bayes accuracy: 72.91

F1-Score: 66.92

Cohen Kappa score: 45.04

Support Vector Machine (SVM)

```
In [29]: from sklearn.svm import SVC

svc_model = SVC(kernel = 'rbf', random_state = 0)

svc_model.fit(X_train, y_train)

y_pred = svc_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("SVM accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.85	0.83	0.84	1166
1	0.82	0.84	0.83	1067
accuracy			0.84	2233
macro avg	0.83	0.84	0.84	2233
weighted avg	0.84	0.84	0.84	2233

SVM accuracy: 83.52
F1-Score: 82.95
Cohen Kappa score: 67.01

Decision Trees

```
In [33]: from sklearn.tree import DecisionTreeClassifier

dt_model= DecisionTreeClassifier(criterion='gini', max_depth=10, random_state=0, min_sam

dt_model.fit(X_train, y_train)

y_pred = dt_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("Decision tree accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.85	0.80	0.82	1166
1	0.79	0.84	0.82	1067
accuracy			0.82	2233
macro avg	0.82	0.82	0.82	2233
weighted avg	0.82	0.82	0.82	2233

Decision tree accuracy: 82.04

F1-Score: 81.75

Cohen Kappa score: 64.1

Random Forest

```
In [31]: from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=0)

rf_model.fit(X_train, y_train)

y_pred = rf_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("Random Forest accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.86	0.81	0.84	1166
1	0.81	0.86	0.83	1067
accuracy			0.83	2233
macro avg	0.83	0.83	0.83	2233
weighted avg	0.83	0.83	0.83	2233

Random Forest accuracy: 83.34

F1-Score: 83.08

Cohen Kappa score: 66.7

K-Nearest Neighbors

```
In [35]: from sklearn.neighbors import KNeighborsClassifier

knn_model = KNeighborsClassifier(n_neighbors = 4, algorithm='ball_tree')

knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)

cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels


cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("KNN accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.69	0.88	0.78	1166
1	0.82	0.57	0.67	1067
accuracy			0.73	2233
macro avg	0.76	0.73	0.72	2233
weighted avg	0.75	0.73	0.73	2233

KNN accuracy: 73.44
 F1-Score: 67.29
 Cohen Kappa score: 46.09

Gradient Boosting Classifier

```
In [37]: from sklearn.ensemble import GradientBoostingClassifier

gbc_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.8, max_depth=2,
gbc_model.fit(X_train, y_train)

y_pred = gbc_model.predict(X_test)

confusion_matrix = confusion_matrix(y_test, y_pred)


cm_display = ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, display_labels
cm_display.plot()
plt.show()

print(classification_report(y_test,y_pred))

acc = accuracy_score(y_test,y_pred)*100
print("GBC accuracy: ",acc.round(2))

f1=f1_score(y_test,y_pred)*100
print("F1-Score: ",f1.round(2))

cohen_kappa = cohen_kappa_score(y_test, y_pred)*100
print('Cohen Kappa score: ',cohen_kappa.round(2))
```



	precision	recall	f1-score	support
0	0.84	0.84	0.84	1166
1	0.82	0.83	0.83	1067
accuracy			0.83	2233
macro avg	0.83	0.83	0.83	2233
weighted avg	0.83	0.83	0.83	2233

GBC accuracy: 83.43
 F1-Score: 82.73
 Cohen Kappa score: 66.81

SVM model produced the highest accuracy score of 83.52 but K-Nearest neighbors had the lowest number of false positives at 136.

Depending on the business strategy model, the marketing department could select the most ideal model based on the results of the accuracy, f1 score, and Cohen-Kappa score.

In []: